

1. Weighted fair queueing

(a)

Start time	Service counters			Packet served	Queued packets	Comments
	A	B	C			
0	0	0	0	A1	∅	
8	8	0	0	--	∅	No packets so service counters should be reset
10	0	0	0	A2	B1	Tie so serve in alphanumeric order
18	8	0	0	B1	∅	Smin = 0 for flow B
38	8	5	5	B2	∅	queues for flows A and C empty, Smin = 5
58	8	10	5	C1	A3, B3, C2, C3	
68	8	10	15	A3	A4, B3, C2, C3	
76	16	10	15	B3	A4, C2, C3	
96	16	15	15	C2	A4, C3, C4	Smin = 15
106	16	16	25	A4	A5, C3, C4	Smin = 16, flow B loses credit
114	24	24	25	A5	C3, C4	Smin = 24, flow B empty
122	32	25	25	C3	C4	Smin = 25, flow B loses credit
132	35	35	35	C4	∅	Smin = 35
142	0	0	0	∅	∅	

(b). Comparing the queueing delay for the packets from flow B, only the packet B3 is served at a different time. With fair queueing (and no weights) B3 starts transmission at time 94, but with WFQ (and a larger weight for flow B), packet B3 starts transmission at time 76.

Notice that we could use a different rule to decide which packet to service. Instead of just picking the flow with the lowest value for S , we could use the current value of S plus the value that is added to S if its next packet is served. For example, at time 10 there are packets in the queue from flows A and B. With the basic rule we pick the packet from flow A because there is a tie in the service counters (they are both 0) and we just pick the flows in alphanumeric order. However, if we used the packet size (adjusted by the weights) we would pick B1 (with a value of 5) over A2 (with a value of 8).

In the long run, either approach for selecting packets will allocate the capacity fairly. But, the rule that considers the current amount of service plus the (weighted) size of the next packet will result in lower queueing delays.

For your information, here is the packet service order if we use the service count plus weighted packet size in selecting packets. Notice in this example the queueing delay for flow B's packets is even lower.

Start time	Service counters			Packet served	Queued packets	Comments
	A	B	C			
0	0	0	0	A1	∅	
8	8	0	0	--	∅	

10	0	0	0	B1	A2	Pick B1 (5) over A2 (8)
30	0	5	0	A2	B2	A2 (8), B2 (10), Smin = 0
38	8	5	5	B2	∅	queues for flows A and C empty, Smin = 5
58	8	10	5	B3	A3, C1, C2, C3	A3 (16), B3 (15), C2 (15). Tie so pick alphanumeric
78	8	15	5	C1	A3, A4, C2, C3	A3 (16), C1 (15), Smin = 5
88	8	15	15	A3	A4, C2, C3	A3 (16), C2 (25), Smin = 8
96	16	15	15	A4	C2, C3, C4	A4 (24), C2 (25) Smin = 15
104	24	15	15	C2	A5, C3, C4	Smin = 15, A5(32), C2(25)
114	24	24	25	A5	C3, C4	Smin = 24 A5(32), C3(35)
122	32	25	25	C3	C4	
132	35	35	35	C4	∅	
142	0	0	0	∅	∅	

2. Additive increase – multiplicative decrease.

The following packets are lost: 9, 25, 30, 38, 50.

The window size is initially 1; when we get the first ACK it increases to 2. At the beginning of the second RTT we send packets 2 and 3. When we get their ACKs we increase the window size to 3 and send packets 4, 5 and 6. When these ACKs arrive the window size becomes 4. Now, at the beginning of the fourth RTT, we send packets 7, 8, 9, and 10; by hypothesis packet 9 is lost. So, at the end of the fourth RTT we have a timeout and the window size is reduced to $4/2 = 2$.

We have the following pattern:

RTT	1	2	3	4	5	6	7	8	9
Packets sent	1	2-3	4-6	7-10	9-10	11-13	14-17	18-22	23-28
Window size	1	2	3	4	2	3	4	5	6
Packet lost				9					25
Ack received	1	3	6	8	10	13	17	22	24

RTT	10	11	12	13	14	15	16	17	18	19
Packets sent	25-27	29-32 (28 not resent)	30-31	33-35 (32 not resent)	36-39	38-39	40-42	43-46	47-51	50-51
Window size	3	4	2	3	4	2	3	4	5	2 (or 3 your choice)
Packet lost		30			38				50	
Ack received	28 *	29	32	35	37	39	42	46	49	51

Note in the 10-th RTT, the source sent packets 25-27. The destination already has packets 26, 27, and 28, and is just waiting for 25. When the destination gets packets 26 and 27 it discards them as duplicates. Since the destination has now received all packets through 28, it sends the cumulative ack 28 to the source.

Note in the 11-th RTT, the destination receives packets 31, and 32. It is missing packet 30, so it sends an ack with number 29. For your information, consider a different alternative. If all four packets (29-32) were lost the destination gets nothing in this round trip time and does not send any acknowledgement (the

destination does not know if the source is done sending packets or if packets were lost, so the destination must wait for the source). In this situation, the source must time out and retransmit.

(A very important note: this problem has considered a simplified congestion control algorithm using AIMD only, and ignores time outs. Also, we consider one ACK per round trip time instead of an ack for each packet. TCP is more complicated because it uses AIMD, slow start, potentially a separate ack for each packet, fast retransmit and fast recovery, etc.)

3. Adapting the Congestion Window.

In this problem we consider a simple network and examine the throughput under three scenarios. In the first scenario, you are given the optimal congestion window size and timeout value. The other two scenarios consider adapting the congestion window using slow start and two different mechanisms to trigger retransmissions.

A simplified model for a network is considered. Suppose that between nodes A and B there is a router R. The link from A to R has an infinite bandwidth (i.e., packets are not delayed). However, the link from R to B introduces a transmission delay of 1 second per packet (e.g., if 2 packets arrive at R it takes 1 second to transmit the first packet, and the second packet waits for 1 second before being transmitted during the next second). To simplify the model, assume that acknowledgments from B through R back to A are sent instantaneously. Further assume that router R has queue size of one, in addition to the packet it is sending. At each second, the sender A first processes any arriving ACK's and then responds to any timeouts. Assume the application process at source A has an infinite backlog of packets to send to destination B, and that packets are numbered 1, 2, 3, ...

Assume that the congestion window is fixed at one packet, and the timeout time set for each packet sent is 1 second. Note that at time 0, source A sends 1 packet because its congestion window is fixed at one, the packet is immediately transmitted by the router R (and the router's queue is empty). At time 1 the packet has completely arrived at the destination B, and an ACK is instantaneously returned to the source A. Source A cancels the timer for packet 1 and sends packet 2. Complete the table below. (Hint: this part is trivial, and the purpose is to just get you familiar with the table structure.)

Time	A receives ACK for packet #	Size of A's congestion window	A sends packet #	R sends packet #	R queues packets
0	-	1	1	1	∅
1	1	1	2	2	∅
2	2	1	3	3	∅
3	3	1	4	4	∅
4	4	1	5	5	∅
5	5	1	6	6	∅
6	6	1	7	7	∅
7	7	1	8	8	∅

(Comment on solution: this is just a very simple version of stop-and-wait. Note that this is the best window size for this problem because the link from R to B can only send 1 packet per second.)

As you found above, the network works well if the window size and timeout values are known. The problem is that TCP does not know the best window size or timeout value. For this part of the problem assume the network is the same and that the timeout value is fixed at 2 seconds (i.e., we will not adapt the timeout value). Now A sends data packets to B over a TCP connection, using slow start. The TCP buffer sizes at both A and B are arbitrarily large, so ignore the advertised window size. For the remainder of this

problem we will assume that when TCP encounters a timeout it reverts to stop-and-wait as the outstanding lost packets in the existing window at the time of the timeout get retransmitted one at a time. The slow start phase does not begin again until all the packets in the existing window are acknowledged. Also, once one timeout and retransmission is pending, subsequent timeouts of later packets are suppressed until the earlier acknowledgment is received. That is, make a list of timeouts that occur and process them one at a time.

Fill in the table below. Note that to fill in the table you will need to keep track of some additional information. For example, you will need to keep track of when timers are set and if they expire, the list of packets that are in the window at both the source and destination, which packets are dropped, etc.

Time	A recvs ACK #	Size of A's congestion window	A sends packet #	R sends packet #	R queues packets
0	-	1	1	1	∅
1	1	2 [2, 3]	2, 3	2	3
2	2	3 [3, 4, 5]	4, 5	3	4 (5 lost)
3	3	4 [4, 5, 6, 7]	6, 7	4	6 (7 lost)
4	4 (timeout packet 5)	1	5	6	5
5	4	1	Nothing waiting for ack 5	5	∅
6	6	1	7	7	∅
7	7	1 [8]	8	8	
8	8	2 [9, 10]	9, 10	9	10

Comment: The most important point is that in slow start, each time an ACK for a packet is received the congestion window is increased by 1 packet. This, of course, only applies when the packet is transmitted and not lost. If there is a retransmission then the ack for the retransmission does not count in increasing the congestion window size.

So, at time 2 when the ack is received for packet 2 the congestion window is immediately increased to 3. Packets 3, 4, and 5 are eligible to be transmitted but packet 3 is already outstanding. So, node A transmits packets 4 and 5. Note that packet 5 is lost because the router can only queue 1 packet while it is waiting to transmit packet 3.

At time 3 the source node does not know that there is any problem. It keeps increasing its window and transmitting more packets.

At time 4 the source first notices a problem. Packet 5 was transmitted by the source at time 2, so it expects that ack at time 4. Because there is a timeout at time 4 for packet 5, the source immediately stops using a sliding window and reverts to stop-and-wait. Furthermore, it must use stop-and-wait until all the outstanding packets are acknowledged (that is until packets 5, 6, and 7 are accounted for). So at time 4 the source retransmits packet 5. Note that the router is a layer-3 device so it does not understand TCP headers. Thus, the router does not know that this is a retransmission or just some other packet. So, there is no way for the router to “flush” the queue or cancel packets in transmission. This would violate the layering approach because the router would have to understand TCP and the TCP headers to be able to try anything this complicated.

At time 5, there is a duplicate ack but note that the source cannot transmit another packet. The source is using stop-and-wait so it is waiting for the ack for packet 5 and it cannot transmit another packet until it gets the ack for 5. Thus, the source cannot transmit anything at time 5.

At time 6, the ack for packet 5 is received. Notice that the ack number is 6 because the receiver has already gotten packet 6 and stored it in its buffer. So, it sends one cumulative ack equal to 6 indicating it has received all packets up through 6. Packet 7 is an outstanding packet from before so we next must use stop-and-wait with packet 7.

At time 7, the ack for packet 7 is received and we can now quit stop-and-wait and return to slow start. Slow start always begins with a window size of 1. Note we do not increment the window size when the ack for packet 7 is received because packet 7 was retransmitted. Only increment the window size when a packet is sent with no retransmissions.

At time 8, the congestion window is increased to 2 (and the same problem starts again).

For the above part, you examined how the window size changes and the order packets are transmitted given a fixed timeout. For this part of the problem, we will look at how duplicate ACK's trigger retransmissions instead of timeouts. So, changes from the previous part are:

- The router has buffer space for three packets instead of just one.
- Fast retransmit is done the first time that a second duplicate ACK is received (that is, the third ACK of the same packet). Fast retransmit was not considered in the previous part. Also, ignore fast recovery; when a packet is lost let the window size be one.
- The timeout interval is infinite (that is, ignore all timeout timers as they are not needed for this example).

Time	A recvs ACK #	Size of A's congestion window	A sends packet #	R sends packet #	R queues packets
0	-	1	1	1	$\emptyset, \emptyset, \emptyset$
1	1	2 [2-3]	2, 3	2	3
2	2	3 [3-5]	4, 5	3	4, 5
3	3	4 [4-7]	6, 7	4	5, 6, 7
4	4	5 [5-9]	8, 9	5	6, 7, 8 (9 lost)
5	5	6 [6-11]	10, 11	6	7, 8, 10 (11 lost)
6	6	7 [7-13]	12, 13	7	8, 10, 12 (13 lost)
7	7	8 [8-15]	14, 15	8	10, 12, 14 (15 lost)
8	8	9 [9-17]	16, 17	10	12, 14, 16 (17 lost)
9	8	Send nothing. Window does not change and all packets are outstanding		12	14, 16
10	8	Fast retransmit. Start stop-and-wait for packets 9 through 17	9	14	16, 9
11	8	Still stop-and-wait	Nothing	16	9
12	8	Still waiting	Nothing	9	\emptyset
13	10	Cumulative ack shows 9 and 10 received so 11 is next	11	11	\emptyset
14	12		13	13	\emptyset
15	14		15	15	\emptyset

16	16		17	17	∅
17	17	1 [18] slow start	18	18	∅
18	18	2 [19-20]	19, 20	19	20

4. **TCP trace.** Slow start is active up to about 0.5 sec on startup. At that time a packet is sent that is lost; this loss results in a coarse-grained timeout at $t=1.9$. At that point slow start is again invoked, but this time TCP changes to the linear-increase phase of congestion avoidance before the congestion window gets large enough to trigger losses. The exact transition time is difficult to see in the diagram; it occurs sometime around $t=2.4$. At $t=5.3$ another packet is sent that is lost. This time the loss is detected at $t=5.5$ by fast retransmit; this TCP feature is the one not present in Figure 6.11 of the text, as all lost packets there result in timeouts. Because the congestion window size then drops to 1, we can infer that fast recovery was not in effect; instead, slow start opens the congestion window to half its previous value and then linear increase takes over. The transition between these two phases is shown more sharply here, at $t = 5.7$.

5. **RED.** Chapter 6, number 37.

Only when the *average* queue length exceeds MaxThreshold are packets automatically dropped. If the average queue length is less than MaxThreshold, incoming packets may be queued even if the real queue length becomes larger than MaxThreshold. The router must be able to handle this possibility.

6. **ECN versus RED.**

The explicit congestion notification (ECN) proposal suggests that a source node using TCP use multiplicative decrease when it receives an ACK with the ECN bit set. This suggests that relatively few of the packets should have the ECN bit set. Instead a congested router should set the bit once for a particular flow and then allow that flow enough time to respond to the ECN bit. Note it will take the round trip time (RTT) before the router will experience any change in its traffic due to the flow that has been marked. (The ECN bit has to travel in the marked packet to the packet's destination and the destination host returns the ECN bit in the ACK back to the source host. The source host reduces its congestion window. So, it takes a RTT for the first packets that arrive back at the congesting router after the router has set the ECN bit.)

The mechanism for marking packets used with DEC bit is to mark all packets if the queue size is above some average value. This can result in a large burst of packets having the bit set. If more than one of these packets are from one particular TCP flow, the source of this flow will end up reducing its window size for each marked packet. This can result in a much larger change in the source's congestion window than is necessary.

7. **TCP**

Flow and congestion control do not play a role. Flow control is concerned with the buffer space at the receiver and uses the advertised window to limit the sender. In this problem there is sufficient buffer space at the receiver so the advertised window plays no role. Congestion control can grow the sender's window size but the sender's window size is limited by the sender's buffer size. That is, there is no value in having the sender's window size larger than its buffer as there would be no packets to send. Note that the sender's buffer can only hold 1% of the file.

The ability of the process to pass data to TCP depends on how quickly the ACK's are received to clear the sender's buffer. So, the rate is limited by the rate ACK's are returned. This is called self-clocking.